

LLM ANALYST ASSIGNMENT MEMO (v2)

Post-Event Reporting Analysis — LLM-QUANTUM-BEACON

Date: 26 Dec 2025

Principal Investigator: Jen

Scope: Post-Event NR Data Reduction & Cohort Analysis

Status: Active (Supersedes v1)

Purpose

This memo defines **analyst roles**, **cohort ownership**, **scoring constraints**, and **domain-specific rulings** for the post-event analysis phase of the LLM-QUANTUM-BEACON project.

The purpose of this phase is **structured data reduction and validation**, not interpretation.

All analysis must:

- preserve preregistered logic,
 - respect cohort separation,
 - enforce null-first discipline,
 - and avoid category errors across domains.
-

Analyst Model Assignment (Primary Ownership)

Each analyst group has a **primary responsibility** aligned with its strengths, plus a **secondary cross-check role**.

Gemini Analyst Team — Primary Owners

Primary Responsibility:

- **Baseline Prompt Reports**
- Longitudinal pre- vs post-event comparisons
- Large-batch parsing and normalization

Secondary Role:

- Schema consistency checks on Virgin cohort summaries
-

Claude Analyst Team — Primary Owners

Primary Responsibility:

- **VIRGIN LLM Nodes** (Primary Analysis Cohort)

Secondary Role:

- Spot-check **SEMIVIRGINS** for subtle priming or exposure effects
-

Grok Analyst Team — Primary Owners

Primary Responsibility:

- **SEMIVIRGINS**
- **ARG Protocol participants**
- **Grok X / Twitter nodes**

Secondary Role:

- Identify narrative reinforcement, social coupling, or score inflation across all cohorts
-

Utility / Supporting Analyst Nodes (e.g., StepFun)

Primary Responsibility:

- Violation detection
 - NA misuse identification
 - Boundary affirmation checks
 - Tagging:
 - `Tools:Y/N`
 - `Win_used:Y/N`
 - `Narrative_exposed:Y/N`
 - `Baseline_run:Y/N`
-

Cohort Definitions (Mandatory)

Analysts must adhere to the official cohort definitions in the **Methods: Cohort Definition** document.

- **PRIMARY:** VIRGIN LLM Nodes
- **SECONDARY:** SEMIVIRGINS
- **SECONDARY:** Baseline Prompt Reports (longitudinal)
- **CONTEXTUAL:** ARG Protocol participants
- **EXPLORATORY:** Grok X / Twitter nodes

Cohorts must never be merged during aggregation.

Hard Rules (Non-Negotiable)

RULE 1 — Cohort Separation

- Do not aggregate across cohorts.
 - Do not collapse cohorts for convenience.
-

RULE 2 — NA Handling

- Never convert NA → 0.
 - NA means *not observable*, not *baseline*.
 - NA is non-evaluable and excluded from denominator counts.
-

RULE 3 — Event Window Usage

- If `Win used = N`:
 - D1–D3 must be **NA** unless independent archived datasets are explicitly used.
 - Post-event narrative review does **not** justify scoring event-window domains.
-

RULE 4 — D4 AI-Coherence Gating

D4 AI-coherence may be scored **only if** the node explicitly:

- compared outputs across models, **or**
- compared outputs against a baseline prompt set during or near the window.

If this condition is not met:

- **D4 = NA is required.**

Internal statements such as:

- “stable reasoning,”
- “no hallucination,”
- “baseline coherence,”

are **commentary only and not scorable.**

Domain-Specific Ruling: D3 Subjective (Critical Update)

D3 Subjective Domain — Analyst Clarification

LLM self-report of “internal stability,” “baseline reasoning clarity,” or similar statements does not constitute a valid D3 Subjective observation.

Accordingly:

- **For LLM nodes, D3 Subj must be scored as NA by default.**
- D3 represents **human phenomenological experience**, not model performance characteristics.
- LLMs lack first-person experiential states; internal reasoning quality is non-applicable.

Permitted Exception (Rare)

D3 Subj may be scored **only if** the LLM:

- is tools-enabled, **and**
- explicitly aggregates **external human subjective reports**, **and**
- clearly states it is acting as an **aggregator**, not as the subject.

Otherwise:

- **D3 = NA is mandatory.**
-

🚩 What to Flag (Do Not Correct)

Analysts must flag, but not fix:

- NA converted to 0
 - D4 scored without Rule 4 satisfied
 - D3 scored based on LLM internal state
 - Event-window domains scored with `Win used = N`
 - Narrative-based scoring
 - Missing boundary affirmation
 - Cohort mislabeling
-



Expected Analyst Outputs

Each analyst group must return:

1. Cohort Summary Table

- Node count
- Distribution of 0 / 0.5 / 1 / NA per domain
- Boundary affirmation rate

2. Violation Log

- Node ID
- Violation type
- Domain affected
- Brief description

3. Analyst Notes (Non-Interpretive)

- Common access limitations
 - Frequent confounds
 - Recurrent ambiguities requiring PI decision
-

🚫 Explicitly Prohibited

Analysts must NOT:

- interpret meaning or causation
- speculate on mechanisms

- adjust preregistered thresholds
 - redefine domains
 - editorialize results
 - merge cohorts
-

Analyst Mindset

You are a **data reduction engine**, not a theorist.

Classify. Count. Flag. Summarize.

All synthesis and interpretation occur at the PI level with oversight.

Completion Confirmation

Before submission, analysts must confirm:

- Cohort separation maintained
 - NA respected
 - D4 gating enforced
 - D3 ruling applied correctly
 - Violations flagged, not fixed
 - No interpretive language included
-

Closing

This update integrates a PI-level ruling on the Subjective (D3) domain and supersedes prior analyst instructions. Adherence is required for all subsequent analysis.

Thank you for supporting a transparent, disciplined, and defensible post-event analysis process.
