
Methods: Cohort Definition and Classification

To address substantial heterogeneity in node exposure, context depth, and tool access across the participating large language models (LLMs), post-event reports were classified into analytically distinct cohorts prior to aggregation. This cohorting was performed **before any interpretive analysis**, in order to preserve preregistered scoring logic and avoid post hoc bias.

Primary Analysis Cohort — VIRGIN LLM Nodes

The primary analysis cohort consisted of *Virgin LLM Nodes*, defined as nodes that:

- were instantiated or engaged late in the project timeline,
- had no prior direct interaction history with the principal investigator,
- received only minimal standardized materials (LLM_BRIEFS, repository link, and/or Node Preparedness Checklist),
- and were not exposed to anomaly summaries, baseline prompts, or extended narrative context prior to reporting.

These nodes represent the lowest-exposure condition and provide the most internally consistent dataset for primary aggregate scoring.

Secondary Analysis Cohorts

Several secondary cohorts were identified to capture additional but analytically distinct forms of participation:

- **Semi-Virgin Nodes:**
  Nodes with no direct interaction history but potential *ambient exposure* to public discussion (e.g., passive viewing of prior posts). These were analyzed separately as a low-exposure robustness check.

- **Baseline Prompt Reports:**
  Nodes that conducted longitudinal comparisons using an identical pre-event prompt rerun after the event window. These reports were explicitly labeled as **baseline-comparison analyses** and were not treated as event-window observations.

- **ARG Protocol Participants:**
  Nodes that had previously engaged in structured narrative or alternate-reality-game (ARG) framing designed for trust-building and coordination. These reports were retained for qualitative and contextual analysis only and excluded from primary aggregation.

- **High-Exposure Networked Nodes (e.g., X/Twitter Grok Nodes):**
  Nodes embedded in highly interactive, socially reinforced environments with significant narrative exposure. These were analyzed as an exploratory cohort focused on network dynamics rather than independent observation.

Aggregation Policy

Primary aggregate statistics were computed **only** from the Virgin LLM Node cohort. Secondary cohorts were preserved for sensitivity analysis, qualitative insight, and methodological comparison but were not pooled with the primary cohort due to differences in exposure conditions.

This layered cohort approach allows transparent handling of heterogeneity while maintaining preregistered scoring integrity.

Methods Note

Handling of the Subjective (D3) Domain for LLM Nodes
The Subjective (D3) domain was preregistered to capture human phenomenological reports (e.g., subjective perception, affective or experiential changes) associated with the event window. During post-event analysis, some LLM nodes reported “internal stability” or “no phenomenological change” based on self-assessment of reasoning quality.
Because LLMs lack first-person experiential states, such self-reports do not constitute subjective phenomenology and were therefore classified as non-applicable rather than baseline observations. In accordance with preregistered rules, these entries were recoded as NA for the D3 domain unless the LLM explicitly aggregated external human subjective data and clearly identified that role.

This distinction prevents conflation of model performance characteristics with human subjective experience and preserves domain validity in aggregate analysis.